

NLP Assignment 3: Transition Parsing with Neural Networks

Utkarsh Garg

November 16, 2019

1 Transition Parsing System

We use the Arc Standard system to implement the Dependency Parser, which involves 3 components: A stack s , a Buffer b and a set of dependency arcs A .

The initial configuration for a sentence w_1, \dots, w_n is $s = [\text{ROOT}]$; $b = [w_1, \dots, w_n]$, $A = \phi$. A configuration c is terminal, if the buffer is empty and the stack contains the single node ROOT, and the parse tree is given by A_c .

2 Implementation

2.1 Dependency Parser: Arc-Standard Algorithm:

- We start by Implementing the Arc-Standard algorithm in parsing_system.py file. This involves creating 3 transitions viz. Left-Arc, Right-Arc and Shift.
- We're given transitions as inputs in the format $L(\text{label})$ or $R(\text{label})$ in case of Arc operations, otherwise simply S for Shift operation.
- For Arc transitions, we identify the type of arc operation, plus the label from the transition input. We also extract the 1st two words from the existing stack in order to add an arc with label, and eventually pop one item from the stack, making one word-token a dependent on another.
- For **Shift** operations, we simply push the next input token w_i into the stack.
- For **Left Arc** transition, we add an arc $s_0 \rightarrow s_1$ with label and remove 2^{nd} top element s_1 from the stack. Whereas, for **Right Arc** transitions, we add an arc $s_1 \rightarrow s_0$ with label and remove stack top s_0 from the stack.

2.2 Features extraction:

We follow the section 3.1 from the paper **The choice of $S_w; S_t; S_l$** in order to extract our feature-token ids.

- We are given a Vocabulary and a Configuration file with defined modules for performing operations and retrieving words/ids for Stack and Buffer operations.
- We select 18 words from our Stack totally, in the following order: (1) The top 3 words on the stack and buffer: $s_1; s_2; s_3; b_1; b_2; b_3$; (2) The first and second leftmost / rightmost children of the top two words on the stack: $lc_1(s_i); rc_1(s_i); lc_2(s_i); rc_2(s_i), i = 1; 2$. (3) The leftmost of leftmost / rightmost of rightmost children of the top two words on the stack: $lc_1(lc_1(s_i)); rc_1(rc_1(s_i)), i = 1; 2$.

- We use the corresponding POS tags for S_t ($n_t = 18$), and the corresponding arc labels of words excluding those 6 words on the stack/buffer for S_l ($n_l = 12$).

2.3 Neural Network Architecture:

We start by defining our Dependency Parser, which makes use of a classifier powered by a neural network. The neural network accepts distributed representation inputs: dense, continuous representations of words, their part of speech tags, and the labels which connect words in a partial dependency parse.

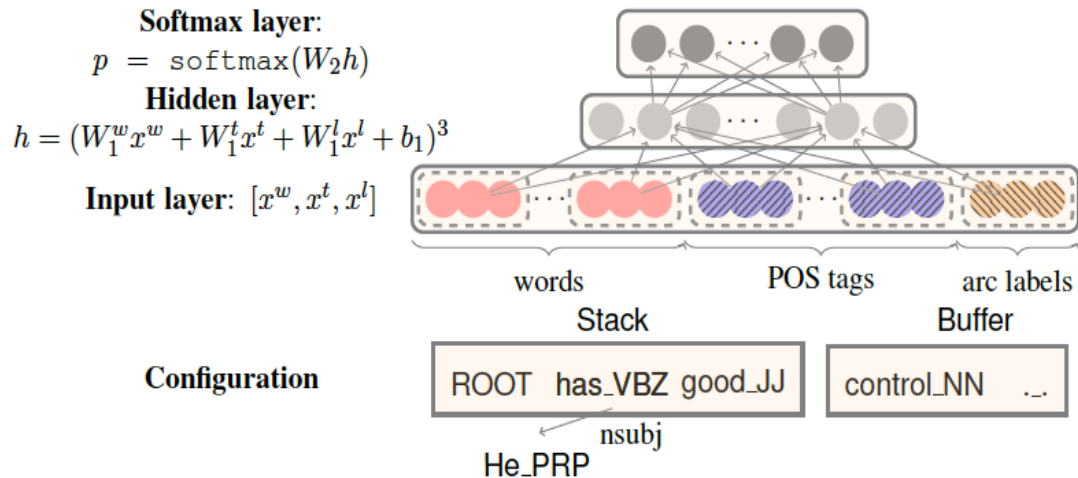


Figure 2: Our neural network architecture.

- In our init call to our Model, we set our activation functions, along with other Trainable variables. We also initialise our embeddings and weights and biases for our subsequent hidden and output layers.
- For our embeddings, we declare a tensor of shape [vocab_size, embedding_dim] with truncated_normal values in the range [-.01, .01].
- For weights and biases for the hidden layer, we do the similar work, this time with a standard deviation defined as $\sqrt{2/fan_in}$, where fan_in represents no. of input features.
- Similarly for the output (classification) layer, we define another weight matrix to be multiplied by output value from the hidden layer.
- In the method call, we then define our Forward Pass for the dependency Parser. We extract the tokens from the embedding matrix initialised earlier, followed by matrix multiplication of hidden layer weights from *weights1* adding the bias.
- This is further converted to non-linear form by using Cubic activation function defined earlier.
- The output from the hidden layer is then subjected to classification layer, using the *weights2* matrix initialised earlier.
- This gives us the logits which are then passed to the compute_loss function for further processing.

2.4 Loss Function:

- For the loss function, we make use of the logits computed from Forward pass implementation, with an aim to obtain the probability of all transitions, incl. feasible and infeasible.
- We define these transitions as follows:
 - 1: 'correct' transition according to the oracle. There is only one 'correct' move per configuration state
 - 0: feasible, but 'incorrect' transition according to the oracle. A move that can be made according to the current state, but would be 'incorrect'
 - -1: infeasible transition. This means that the configuration is not in a state where this transition could be performed.

Since the logits returned from our Model class, represent 3 types of transitions, we need to filter out [-1] transitions in order to compute the correct softmax probabilities.

- For this purpose we create a mask to filter out incorrect transitions/labels [-1] from our computed logits. We have used `tf.cast` on condition `labels_i=0` to get the corresponding float values as 0 and 1 for our mask, identifying infeasible and feasible labels respectively. The result is multiplied with logits to obtain 'feasible_logits'.

- we compute $\exp^{\text{feasible_logits}}$ to compute the numerator for the softmax.

- Since, taking exponent converts labels 0[infeasible] from e^0 value to 1. We therefore need to filter those out. So we reuse the created 'label_mask' to filter only correct $\exp^{\text{feasible_logits}}$. We thus obtain 'masked_exp_feasible_logits'

- Similarly, for the denominator, we compute 'sum_masked_exp_feasible_logits' by taking summation of 'masked_exp_feasible_logits' across y-axis to get a $[\text{batch_size}, 1]$ shaped tensor. We then perform a matrix multiplication with 1's tensor of shape $[1, \text{num_transitions}]$ to obtain the 'sum_masked_exp_feasible_logits' of shape $(\text{batch_size}, \text{num_transitions})$

- We then use the below formula to obtain our final softmax logits(probabilities):

$$\text{softmax}(x) = \frac{\exp x_i}{\sum_i \exp x_i}$$

- The final objective is to minimize the cross-entropy loss, plus a l2-regularization term, given by:

$$L(\theta) = - \sum_i \log p_{t_i} + \frac{\lambda}{2} \|\theta\|^2$$

Here:

$p_{t_i} \rightarrow$ softmax probability for input i.

$\lambda \rightarrow$ regularization_lambda

$\theta \rightarrow$ set of parameters $\{W_1^w, W_1^t, W_1^l, b_1, W_2, E^w, E^t, E^l\}$

- Following above loss function, we compute our log_prob by taking log on our softmax_logits.
- We create a mask to filter just the correct transitions/labels [+1] from our computed $\log(\text{softmax_logits})$, to compute 'correct_logits'.

- We then sum this over entire batch inputs, by taking summation of correct_logits, to get 'batch_loss'.
- Our total loss is computed by taking mean of this batch_loss along axis=1.
- Regularization term is computed by taking sum of L2-loss over set of parameters (*weights1, weights2, bias, embeddings/tokens*), multiplied by regularization_lambda value provided.

3 Analysis

<u>Experiment</u>	<u>UAS</u>	<u>UASnoPunc</u>	<u>LAS</u>	<u>LASnoPunc</u>	<u>UEM</u>	<u>UEMnoPunc</u>	<u>ROOT</u>
Basic (Glove Embeddings) Cubic Activation	88.197023 70566094	89.823659 07421014	85.766632 59964603	87.071158 08511841	35.470588 235294116	38.235294 11764706	90.117647 05882354
Sigmoid	85.771618 0172994	87.605267 6199627	83.301343 57005758	84.807551 00887356	29.882352 94117647	32.352941 17647059	86.764705 88235294
Tanh	87.277214 1486153	88.987170 06725823	84.884213 67500062	86.277058 72378907	34.058823 52941177	36.882352 94117647	87.823529 41176471
Without tunable embeddings	84.944038 68684098	86.715085 06188889	82.274347 53346462	83.694116 31718759	28.0	30.235294 117647058	84.529411 76470588
Without Pretrained (Glove) embeddings	82.271854 82463793	84.457129 93839372	79.836478 30096966	81.778104 33504776	23.411764 70588235	25.588235 29411765	75.529411 76470588

Fig:1 : Performance Evaluation for different Experiments.

We can see that our best model is the Basic one with Pre-trained GLOVE Embeddings , activated by Cubic function. We obtain the ROOT accuracy of 90% , the highest amongst all our experiments. Tanh model is the next closest and has comparable accuracy to Basic.

Going with the best model ,we have thus used the **Basic** Model to generate predictions for the test data (inside test_predictions.conll).